

Class 3 – Data Visualization

Visual Perception & Cognitive Load

Data Types, Encodings & Grammar of Graphics

Python Visualization Fundamentals

Class Overview

Today's Learning Objectives

- Understand how humans perceive visual information
- Learn about cognitive load and its impact on visualization design
- Master data types and their valid visual encodings
- Explore the Grammar of Graphics framework
- Implement visualizations using Python (Matplotlib, Seaborn)
- Apply perceptual principles to create effective visualizations

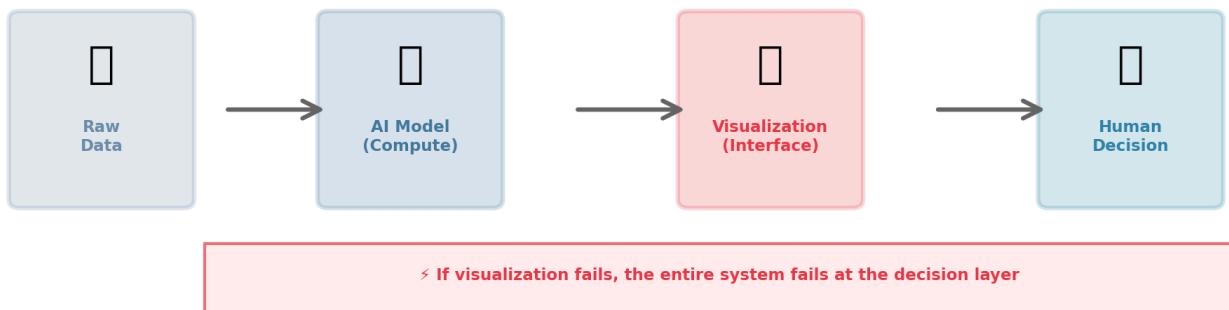
Course Context: Building on Week 2's Python introduction

Why Visualization Matters in Data Science

Visualization is not just presentation—it's cognition

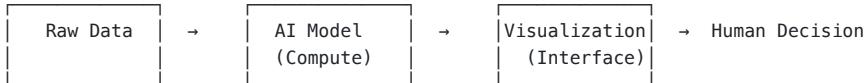
- Data visualization is the interface between:
 - Computation (algorithms, models)
 - Human decision-making
- Every visualization choice affects:
 - Speed of insight
 - Accuracy of interpretation
 - Quality of decisions
- Poor visualization = Poor analytics outcomes

Data to Decision Pipeline: Visualization as Interface



Visualization as a Human–AI Interface

In modern analytics and AI systems:



- Models compute predictions
- Humans make decisions
- **Visualization is the critical translation layer**
- Human perceptual limits become system constraints

Key Insight: If the visualization fails, the entire system fails at the decision layer.

The Two Pathways of Visual Processing

Visual processing happens in two stages:

1. Early Vision (Preattentive Processing)

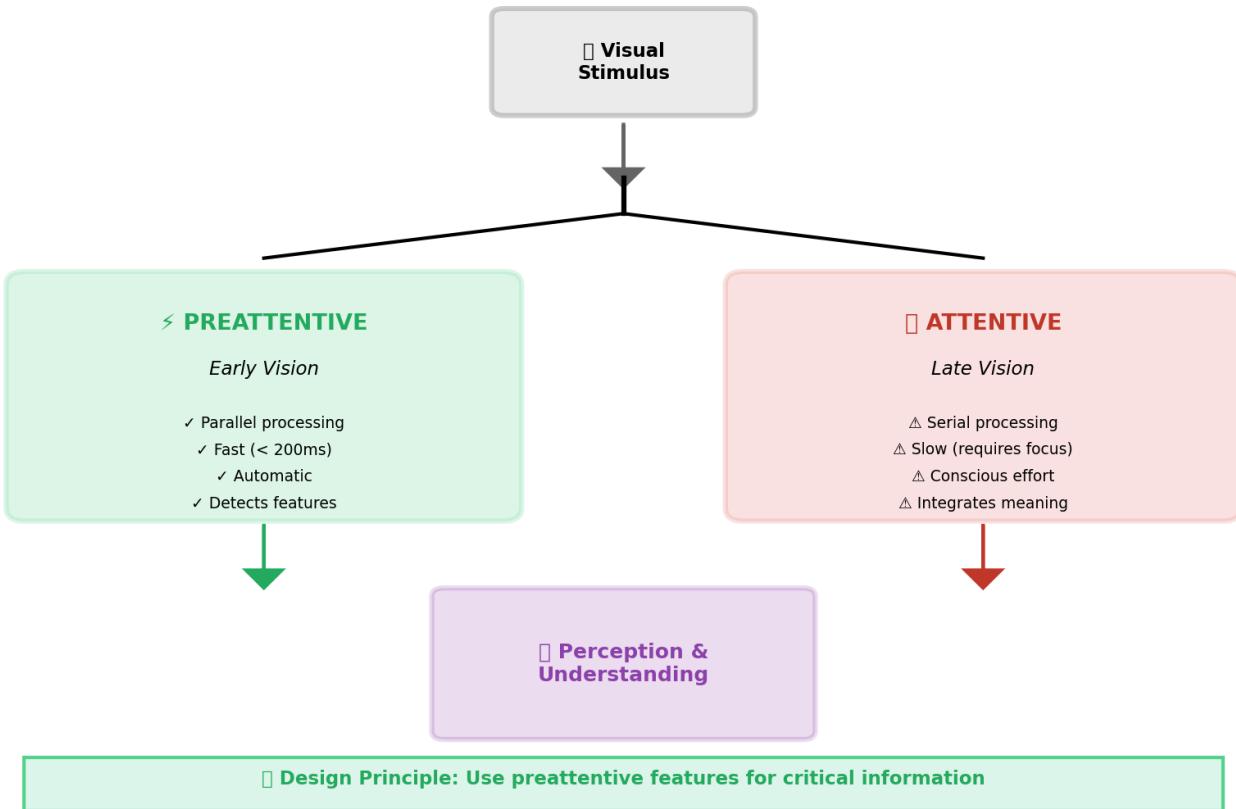
- Parallel processing
- Fast (< 200-250 milliseconds)
- Automatic, no conscious effort
- Detects basic features

2. Late Vision (Attentive Processing)

- Serial processing
- Slow (requires focus)
- Conscious effort required
- Integrates meaning and relationships

Design Principle: Leverage preattentive processing for critical information.

Two Pathways of Visual Processing



Preattentive Processing: The 200ms Window

Certain visual features are detected instantly

Preattentive features enable:

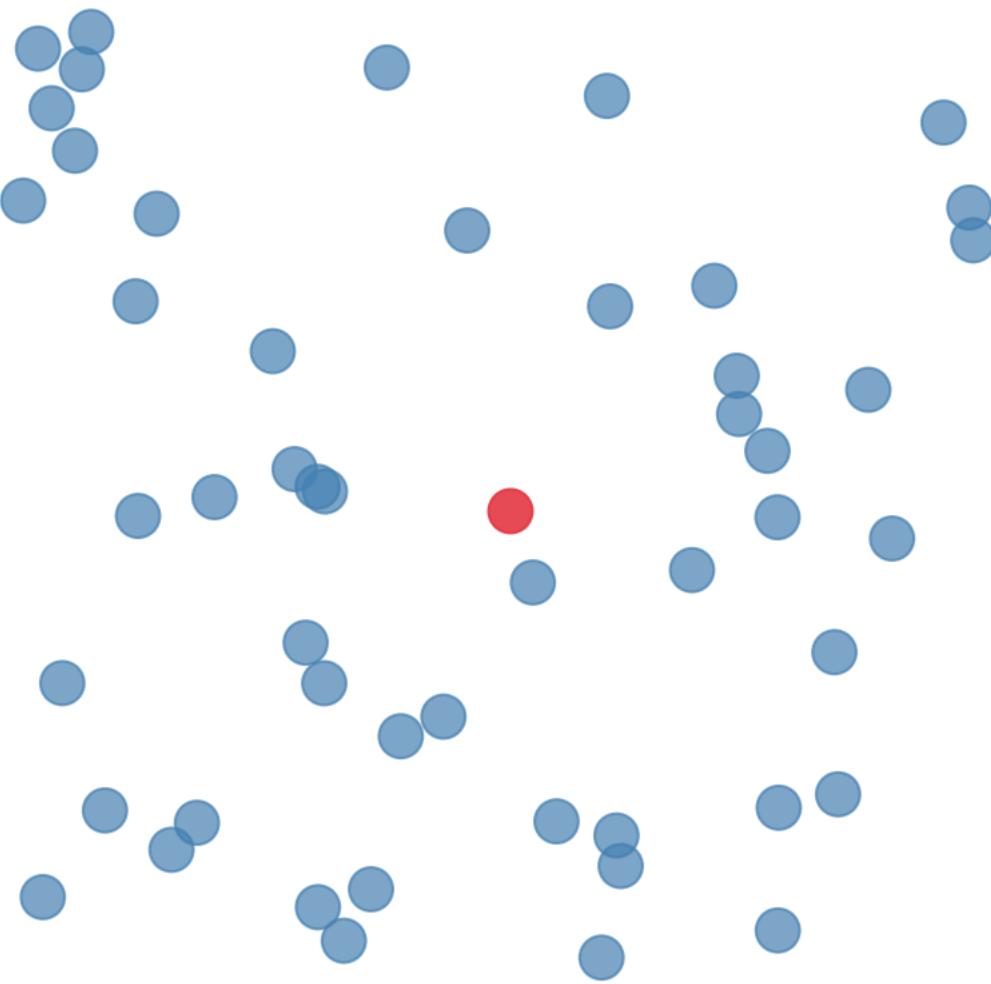
- ⚡ Rapid detection
- ⚡ Visual pop-out effect
- ⚡ Fast filtering and scanning
- ⚡ Parallel comparison

Examples of preattentive features:

- Position
- Color (hue)
- Intensity/brightness
- Size/length
- Orientation/angle
- Shape
- Motion
- Enclosure

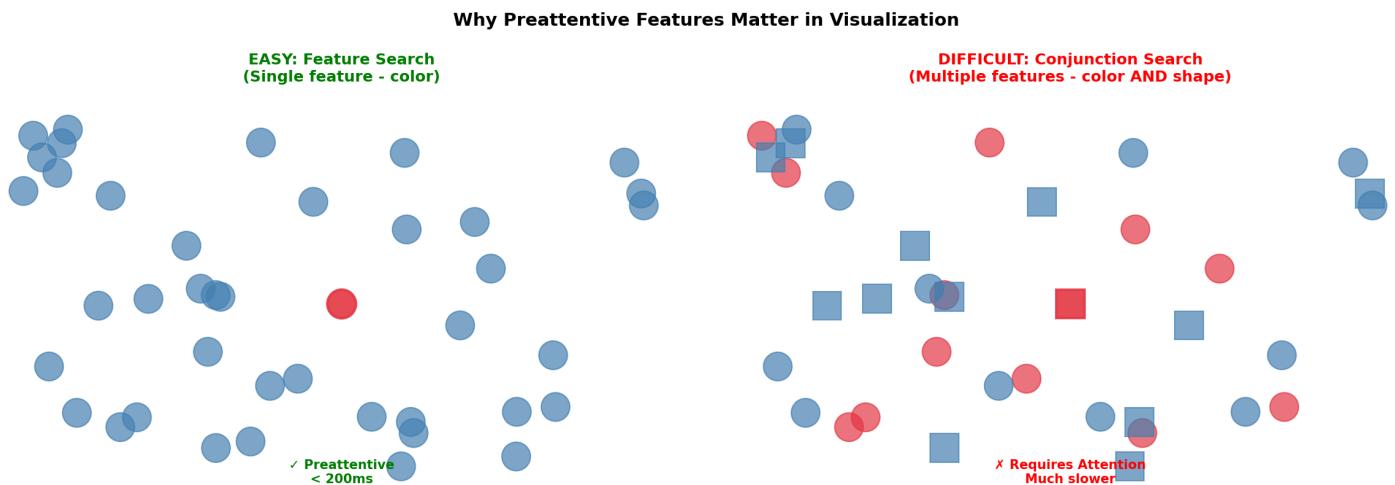
Key Design Principle: Use preattentive features for what matters most in your visualization.

Preattentive Pop-Out: Find the Red Circle



Demonstrating Preattentive Processing

Interactive Exercise: Compare these two search tasks



Key Findings:

- **Left (Feature Search):** Find the red circle
 - Uses only ONE attribute (color)
 - Found instantly (< 200ms)
 - **Preattentive** - no conscious effort required
- **Right (Conjunction Search):** Find the red square
 - Requires TWO attributes (color AND shape)
 - Takes much longer (requires scanning)
 - **Attentive** - conscious effort required

Design Implication: Don't force users to search for multiple attributes simultaneously. In dashboards, make critical information stand out using a single preattentive feature (color, size, or position).

Visual Attention is Limited

The harsh reality of human attention:

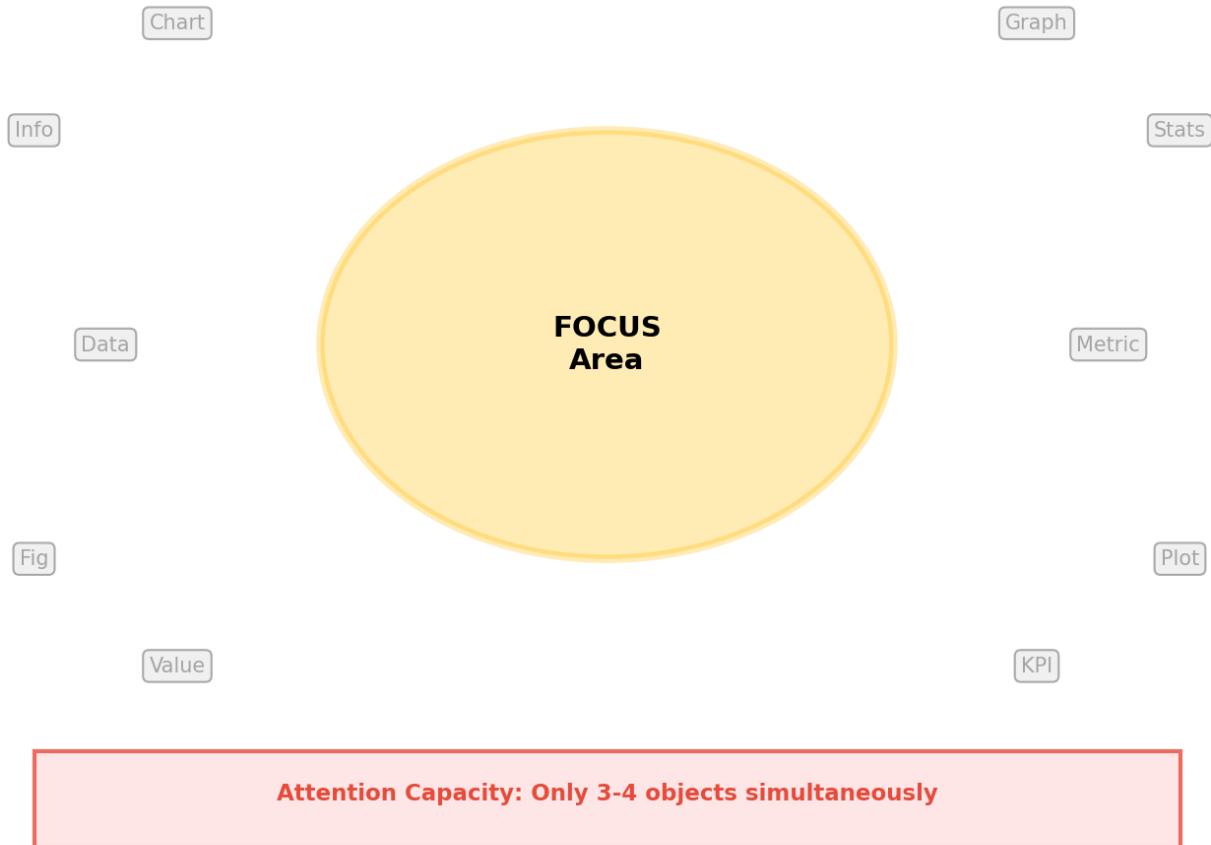
- ✗ Humans cannot attend to everything at once
- ✗ Attention is selective and focused
- ✗ Users must actively search, filter, and compare

Implications for dashboard design:

- If everything is highlighted, nothing is highlighted
- Overloading visual channels creates cognitive fatigue
- Priority must be encoded visually

Attention capacity: ~3-4 visual objects simultaneously

Visual Attention is Limited



Change Blindness

Large changes can go unnoticed if attention is not directed

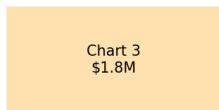
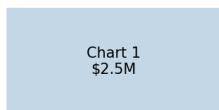
Common in:

- 🖥️ Real-time dashboards (values update but users don't notice)
- 📊 Multi-panel layouts (change in one panel while viewing another)
- 🎥 Animated transitions (too fast or too subtle)
- 📊 Information-dense interfaces

Design Solutions:

- Use motion for important changes
- Add explicit alerts/notifications
- Use color + position changes together
- Minimize unnecessary animation

Dashboard - Before



Dashboard - After



Without directed attention, users often miss significant changes

Inattentional Blindness: The Invisible Gorilla

Famous experiment: Watch a video of people passing basketballs

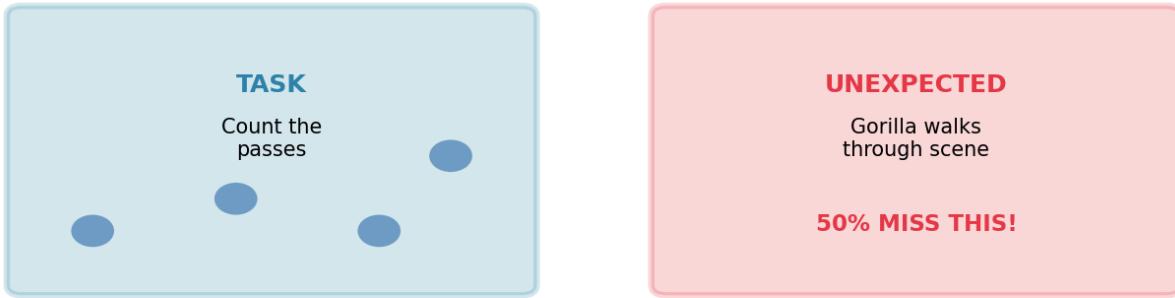
- Task: Count the passes
- Result: 50% of viewers miss a person in a gorilla suit walking through!

Lesson for visualization:

- Users focused on one task will miss other information
- Don't assume users will "just notice" important patterns
- **Make critical information impossible to miss**

Design Principle: If it's important, make it preattentive or explicitly direct attention to it.

Inattentional Blindness



Lesson for Visualization:

Users focused on one task will miss other information

Make critical information impossible to miss

Gestalt Principles: How We Group Information

Gestalt Psychology: The whole is greater than the sum of its parts

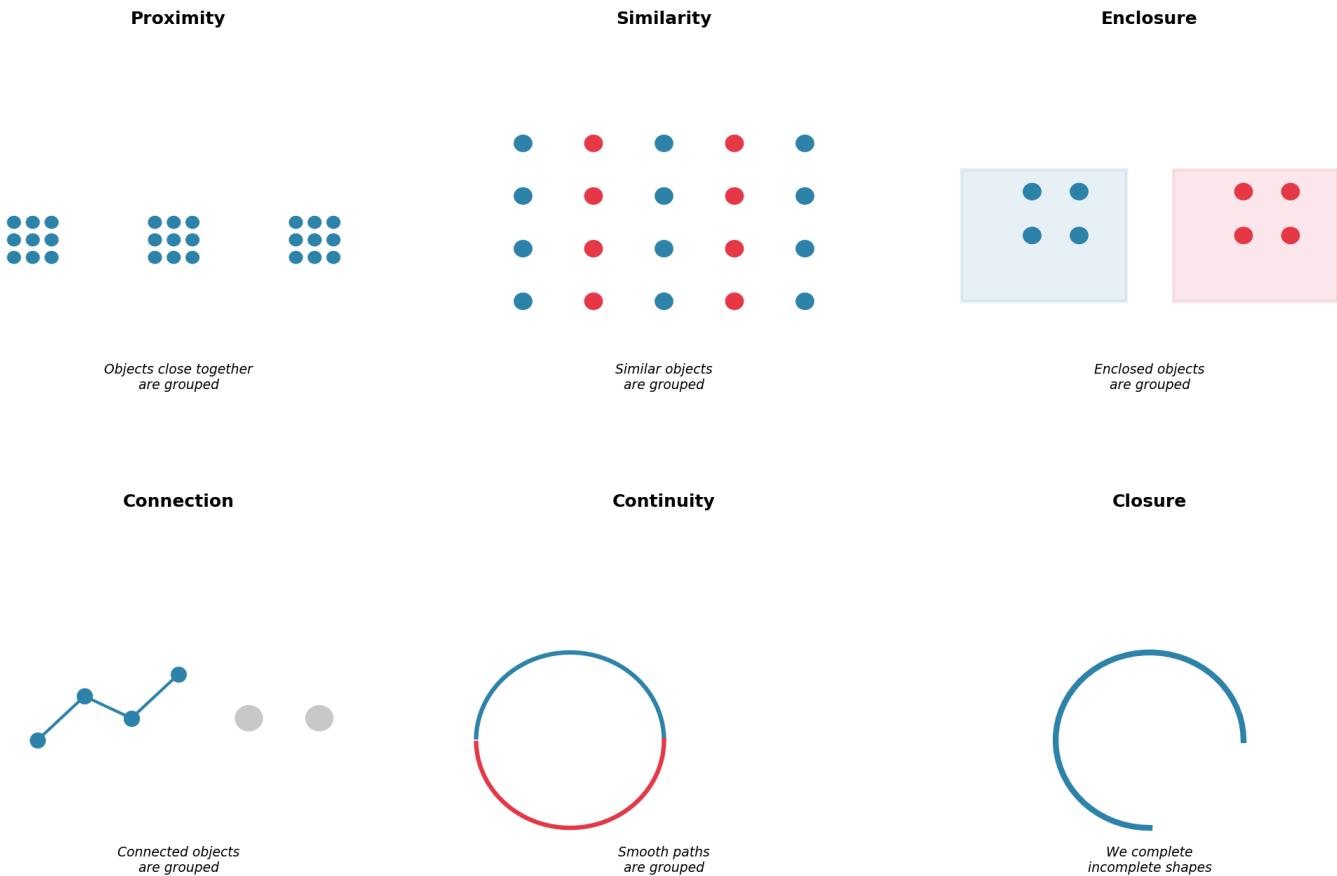
Core principle: Our brains automatically organize visual elements into groups

Key Gestalt principles for visualization:

1. Proximity
2. Similarity
3. Enclosure/Common Region
4. Continuity
5. Connection
6. Closure

Why this matters: Perception overrides logic. Users will group elements based on visual properties, regardless of what your legend says.

Gestalt Principles of Grouping



Gestalt Principle: Proximity

Objects close together are perceived as related

A A A	B B B	C C C
A A A	B B B	C C C
A A A	B B B	C C C

Users will see 3 groups, not 27 individual items.

In visualization:

- Place related charts near each other
- Use whitespace to separate unrelated elements
- Don't rely only on color to show groups
- Accidental proximity creates false relationships

Real-world mistake: Dashboard elements randomly placed create confusing relationships.

Gestalt Principle: Proximity



Objects close together are perceived as grouped

Gestalt Principle: Similarity

Objects that look similar are perceived as related

Similar:

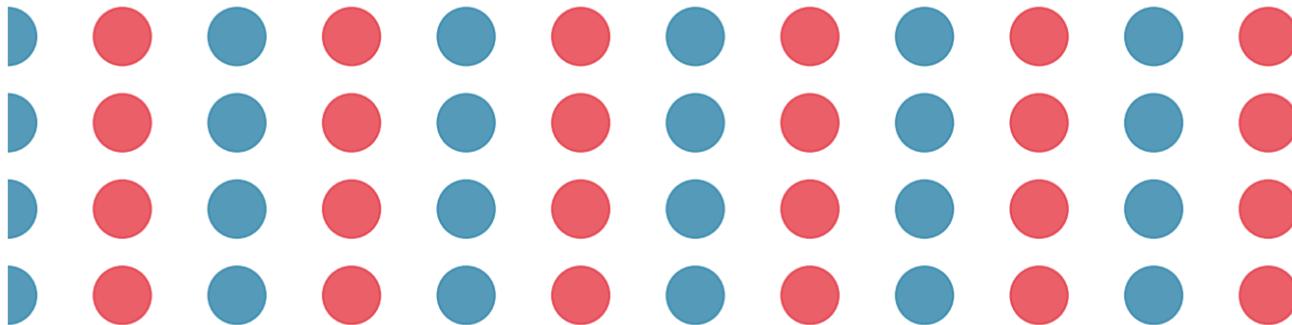
- Color
- Shape
- Size
- Orientation
- Texture

In visualization:

- Use consistent encoding within a group
- Use distinct styles for different data types
- Don't reuse visual encodings inconsistently
- Don't use similar colors for unrelated data

Example: All sales data in blue shades, all costs in red shades.

Gestalt Principle: Similarity



Similar objects are perceived as related

Gestalt Principle: Enclosure & Common Region

Objects within a boundary are perceived as grouped





Power of enclosure:

- Stronger than color alone
- Stronger than proximity
- Creates clear visual hierarchy

In dashboards:

- Use cards/panels for related metrics
- Use background shading for grouped sections
- Be careful with nested enclosures (can create confusion)

Gestalt Principle: Enclosure & Common Region



Enclosure creates strong grouping

Gestalt Principle: Connection

Physically connected elements are perceived as related



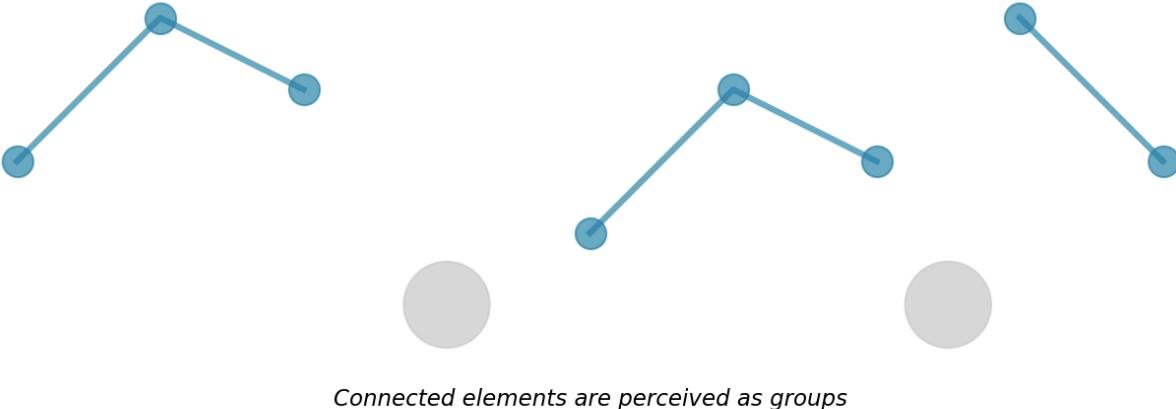
We see: (A-B), C, D, E, (F-G)

In visualization:

- Line charts naturally show relationships
- Connected scatter plots show sequences
- Network diagrams use this principle fundamentally

Design tip: Connection is one of the strongest grouping cues.

Gestalt Principle: Connection



Gestalt and Dashboard Design Mistakes

Common failures when Gestalt principles are ignored:

- Random layout:** Elements scattered without visual hierarchy
- Inconsistent spacing:** Creates accidental groups
- Color chaos:** Similar colors used for unrelated data
- Over-boxing:** Everything has a border, nothing stands out
- No whitespace:** Cognitive overload

Good dashboard design:

- Clear visual hierarchy through proximity and enclosure
- Consistent use of similarity within groups
- Strategic whitespace guides attention
- Visual grouping matches logical relationships

Introduction to Cognitive Load Theory

Core Concept: Working memory has limited capacity

Cognitive Load Theory (Sweller, 1988):

- Human working memory can hold $\sim 7 \pm 2$ chunks of information
- Processing capacity is limited
- Overload leads to errors and decision fatigue

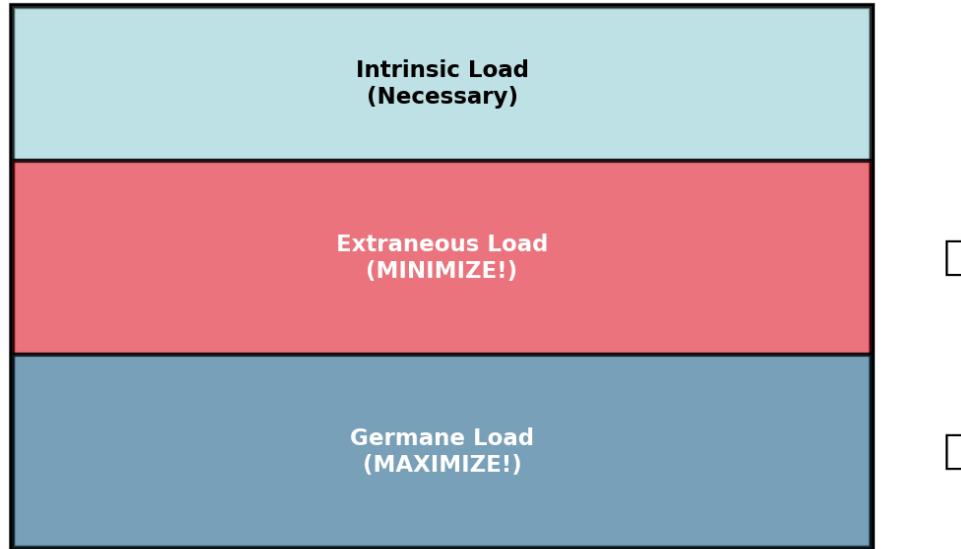
Three types of cognitive load:

1. **Intrinsic Load** - inherent complexity of the task
2. **Extraneous Load** - unnecessary complexity from poor design
3. **Germane Load** - productive mental effort toward learning/insight

Goal in visualization: Minimize extraneous load, manage intrinsic load, maximize germane load.

Cognitive Load Balance

Total Cognitive Capacity (Fixed)



Intrinsic Load: Inherent Complexity

Intrinsic load comes from:

- Data complexity (volume, dimensionality)
- Task complexity (comparison, correlation, prediction)
- Domain complexity (specialized knowledge required)

You cannot eliminate intrinsic load, but you can manage it:

- ✓ **Staging:** Break complex tasks into steps
- ✓ **Scaffolding:** Provide context and guides
- ✓ **Progressive disclosure:** Show detail on demand
- ✓ **Hierarchy:** Present overview first, details later

Example:

- ✗ Showing 50 variables at once → overwhelming
- ✓ Showing top 10, with option to explore more → manageable

Extraneous Load: Unnecessary Complexity

Extraneous load comes from poor design choices:

- ✗ Chart junk (decorative elements)
- ✗ 3D effects with no purpose
- ✗ Cluttered layouts
- ✗ Inconsistent encodings
- ✗ Poor color choices
- ✗ Missing labels or legends
- ✗ Excessive animation

These add NO value but consume cognitive resources

The solution: Ruthlessly eliminate extraneous load

"Perfection is achieved not when there is nothing more to add, but when there is nothing left to take away." — Antoine de Saint-Exupéry

Germane Load: Productive Mental Effort

Germane load supports:

- Understanding patterns
- Making comparisons
- Drawing insights
- Building mental models
- Learning relationships

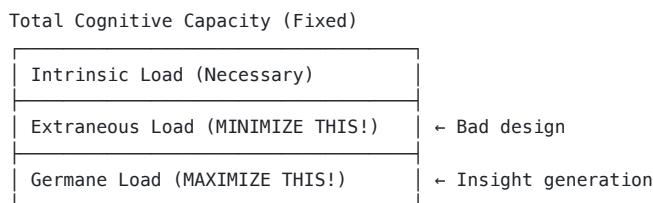
How to maximize germane load:

- Clear visual hierarchy guides thinking
- Comparisons are easy (aligned scales)
- Patterns pop out (preattentive features)
- Context is provided (annotations, benchmarks)
- User can explore (interactive elements)

The goal: All cognitive effort goes toward insight, not decoding the chart.

The Cognitive Load Balance

Visualization design as optimization:



Design principle:

- Every design choice should either reduce extraneous load or increase germane load
- If it does neither, remove it

Real-world impact: Reducing cognitive load = faster decisions = better outcomes

PART 2: DATA TYPES & VISUAL ENCODINGS

Slides 22-33

Cognitive Load in Dashboard Design

Scenario: Executive dashboard with 20 KPIs

X High Extraneous Load Approach:

- All metrics same size and color

- No hierarchy
- 3D pie charts and gratuitous gradients
- Cluttered legends
- Inconsistent time periods

Low Extraneous Load Approach:

- Top 3-5 metrics prominently displayed
- Clear visual hierarchy
- Simple, clean charts
- Direct labeling (no legend hunting)
- Consistent time periods and scales

Result: Same data, drastically different cognitive load and decision quality.

Exercise: Cognitive Load Analysis

Look at two versions of the same dashboard (provided)

For each, identify:

1. Sources of **extraneous load** (what can be removed?)
2. Ways to reduce **intrinsic load** (how to simplify?)
3. Opportunities for **germane load** (what supports insight?)

Discussion: How would redesigning reduce decision time and errors?

Data Types: The Foundation of Encoding

Understanding data types is critical because:

- Data type determines valid encodings
- Violating data type rules creates confusion
- Different data types support different tasks

Four fundamental data types:

1. **Nominal** (Categorical)
2. **Ordinal** (Ordered categories)
3. **Quantitative** (Numerical)
4. **Temporal** (Time-based)

Rule: Data type restricts which visual encodings are appropriate.

Nominal (Categorical) Data

Characteristics:

- Discrete categories
- No inherent order
- Identity and difference matter

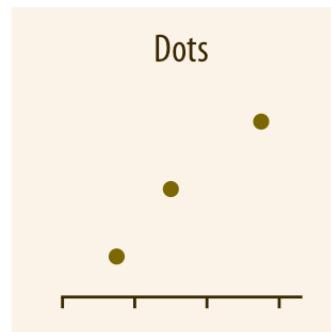
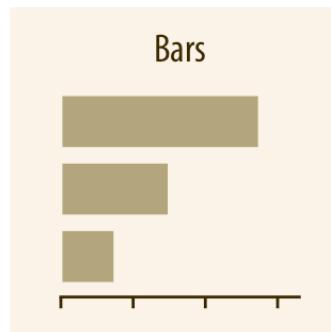
Examples:

- Product categories (Electronics, Clothing, Food)
- Geographic regions (North, South, East, West)
- Customer segments (A, B, C)
- Business units

Valid encodings:

- Hue (color)
- Shape
- Position (grouped)
- NOT size, intensity, or ordering

Common mistake: Using a gradient color scale for categories (e.g., red→yellow→green for product types)



Ordinal Data

Characteristics:

- Discrete categories with order
- Rank matters
- Intervals may not be equal

Examples:

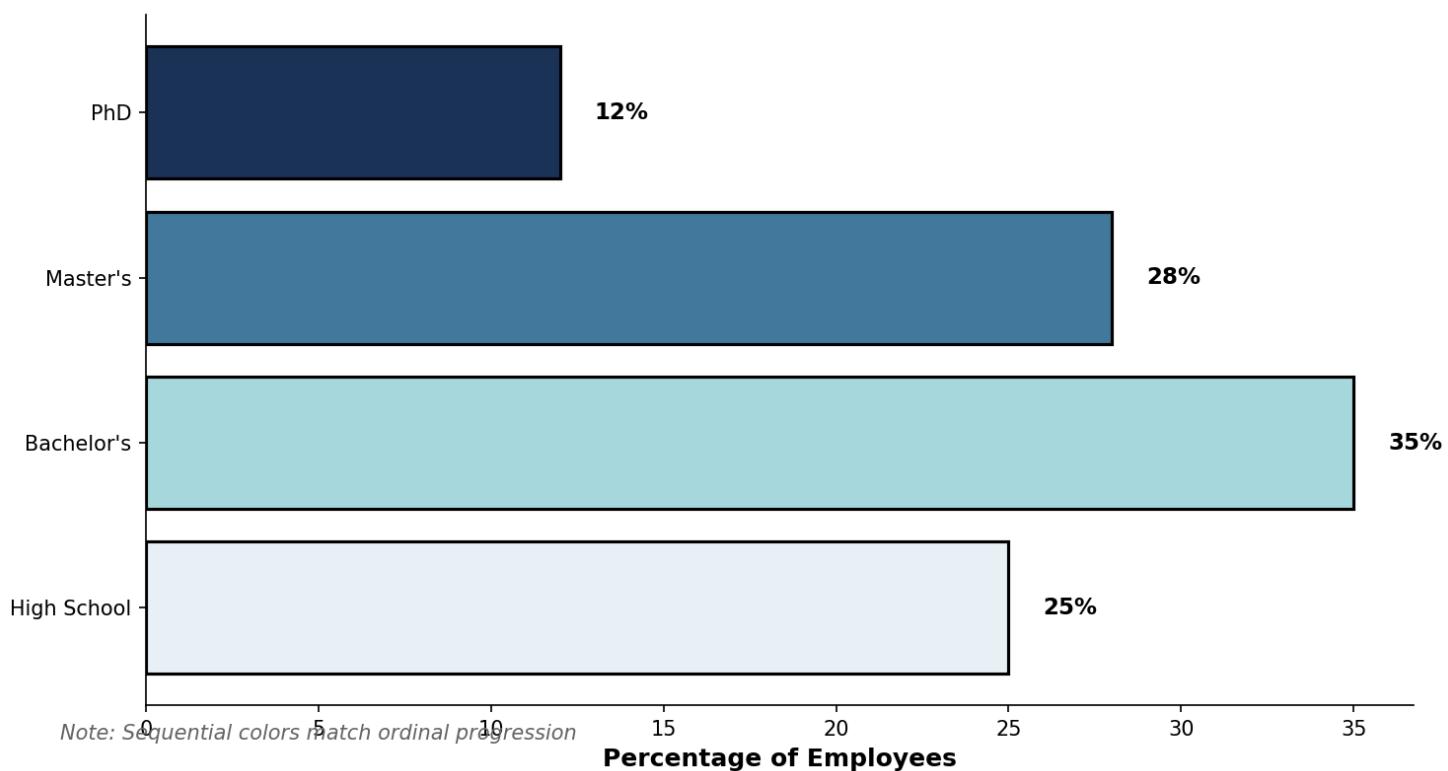
- Education level (High School < Bachelor's < Master's < PhD)
- Customer satisfaction (Poor < Fair < Good < Excellent)
- Priority (Low < Medium < High)
- Size categories (S < M < L < XL)

Valid encodings:

- Sequential color scale (light to dark)
- Position (ordered)
- Size (with caution)
- NOT random colors or shapes

Common mistake: Using arbitrary colors (red, blue, green) for ordered categories.

Ordinal Data: Education Level (Sequential Color Encoding)



Quantitative Data

Characteristics:

- Continuous numerical values
- Magnitude and precise differences matter
- Supports arithmetic operations

Examples:

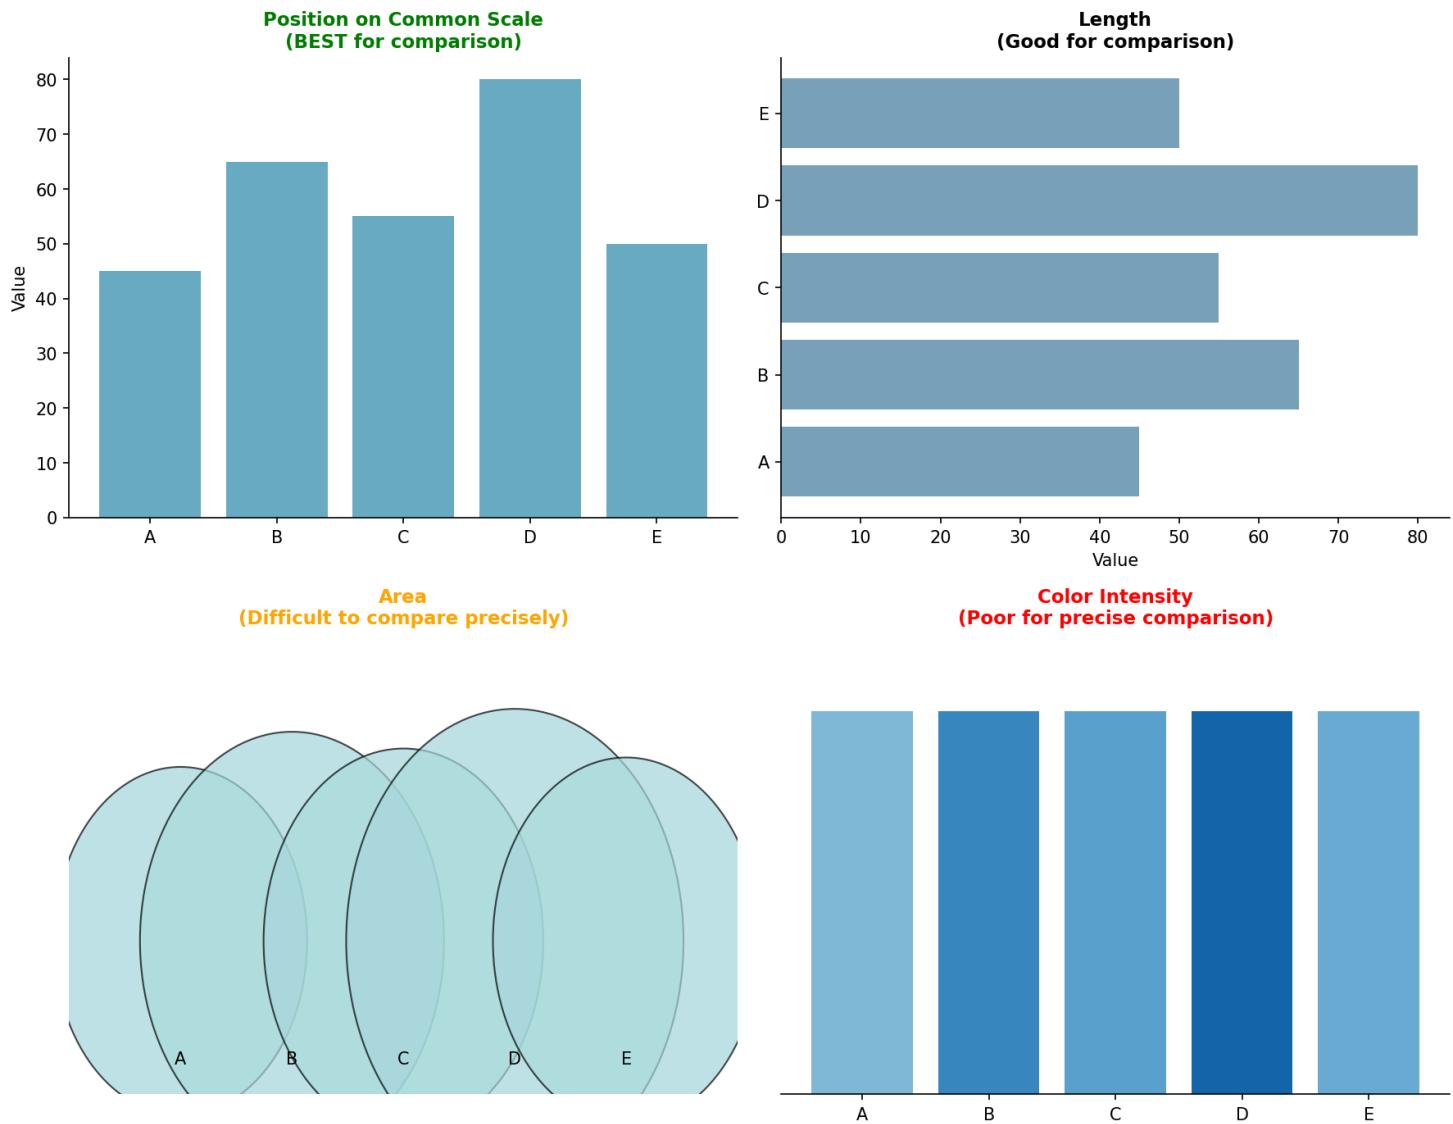
- Revenue (\$1M, \$2.5M, \$3M)
- Temperature (72°F, 85°F)
- Count of customers (100, 250, 500)
- Percentage growth (5%, 12%, -3%)

Valid encodings (ordered by effectiveness):

- Position on common scale (BEST)
- Length
- Area (with caution)
- ⚠ Color intensity (for heatmaps, not precise comparison)
- NOT angle alone, NOT arbitrary hue

Key principle: For precise comparison, use position or length.

Quantitative Data Encodings



Temporal (Time-Series) Data

Characteristics:

- Ordered by time
- Sequential dependency matters
- Trends, cycles, and patterns

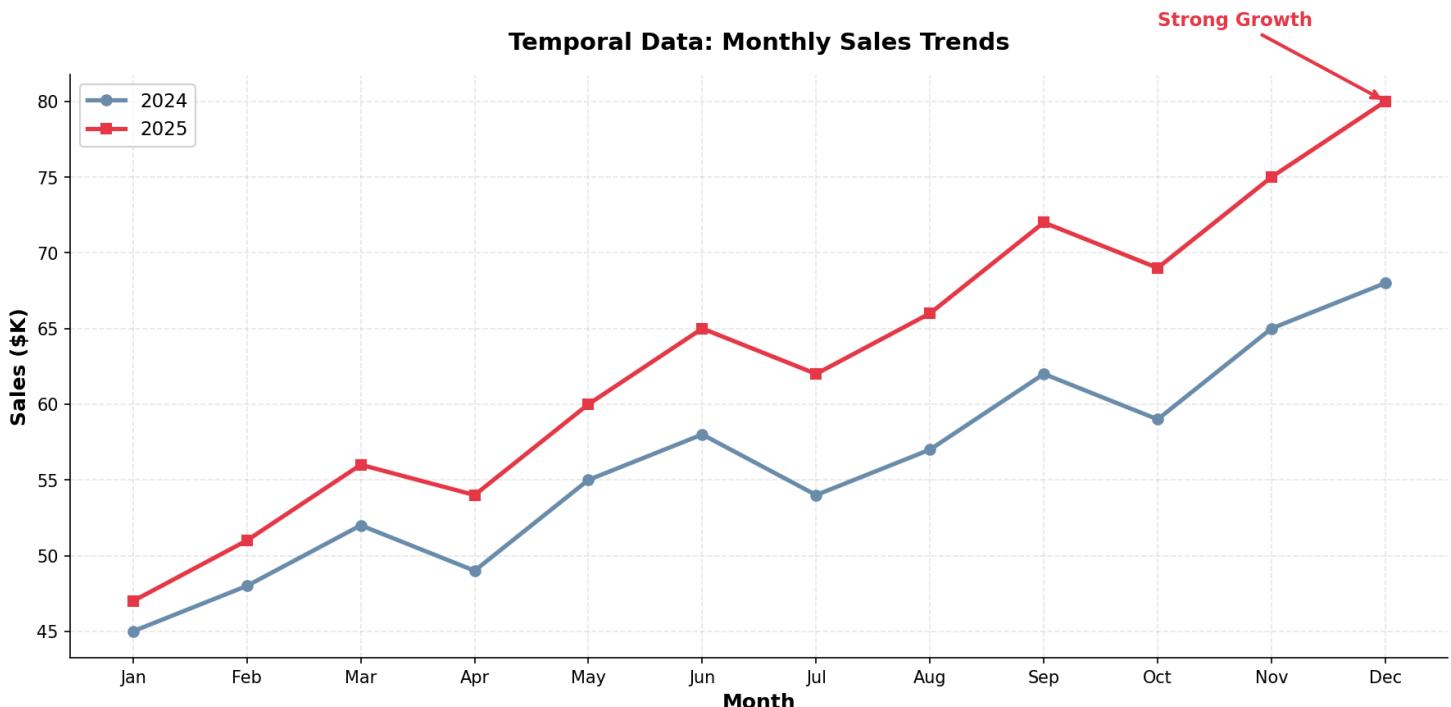
Examples:

- Daily stock prices
- Monthly sales
- Hourly website traffic
- Quarterly earnings

Best practices:

- Time typically on x-axis (horizontal)
- Use line charts for continuous trends
- Maintain consistent time intervals
- Show context (reference periods, benchmarks)
- Don't break time axis arbitrarily

Note: Temporal data is a special case—always respect chronological order!



Visual Variables (Bertin, 1967)

Jacques Bertin identified fundamental visual variables:

1. Position (x, y coordinates)
2. Size (length, area, volume)
3. Value (lightness/darkness)
4. Texture (pattern)
5. Color (hue)
6. Orientation (angle)
7. Shape (form)

Not all variables are equal!

- Some support quantitative judgments (position, size)
- Others support categorical distinctions (hue, shape)
- Some are perceptually weaker (texture, orientation)

Bertin's Visual Variables (1967)

1. Position		<i>x, y coordinates</i>
2. Size		<i>length, area, volume</i>
3. Value		<i>lightness/darkness</i>
4. Texture		<i>pattern, grain</i>
5. Color		<i>hue</i>
6. Orientation		<i>angle, direction</i>
7. Shape		<i>form, symbol</i>

Not all variables are equally effective for all data types

Graphical Perception Research (Cleveland & McGill)

Landmark study (1984): Empirically tested how accurately people decode visual encodings

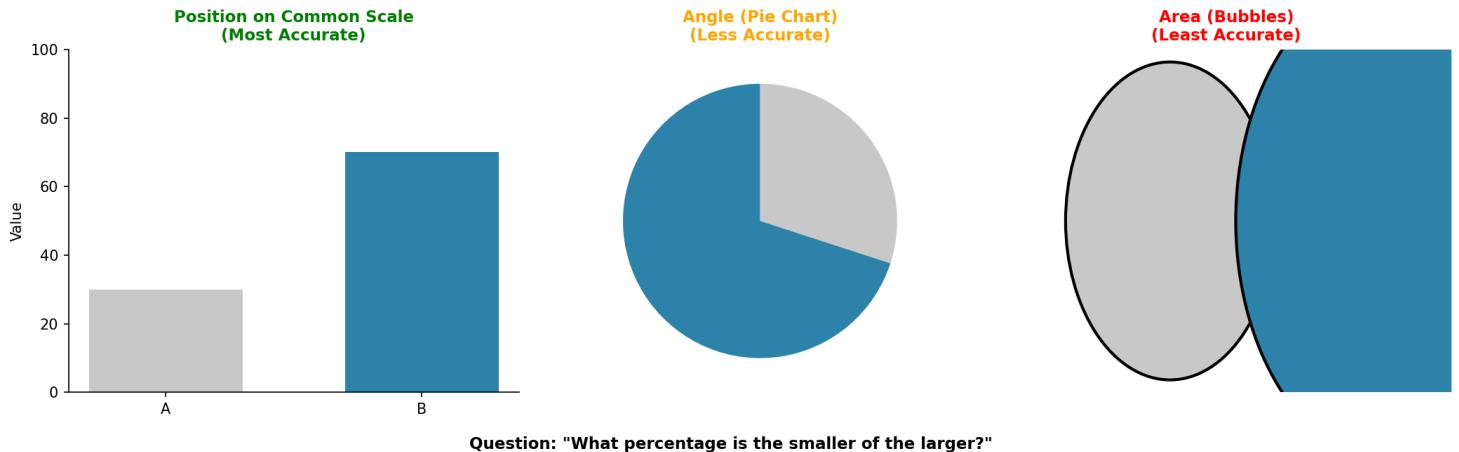
The experiment:

- Show two bars/wedges/circles
- Ask: "What percentage is the smaller of the larger?"
- Measure accuracy and speed

Key finding: Position is dramatically more accurate than angle or area

This research transformed visualization from art to science

Cleveland & McGill's Graphical Perception Study (1984)



Ranking of Visual Encodings (Most→Least Effective)

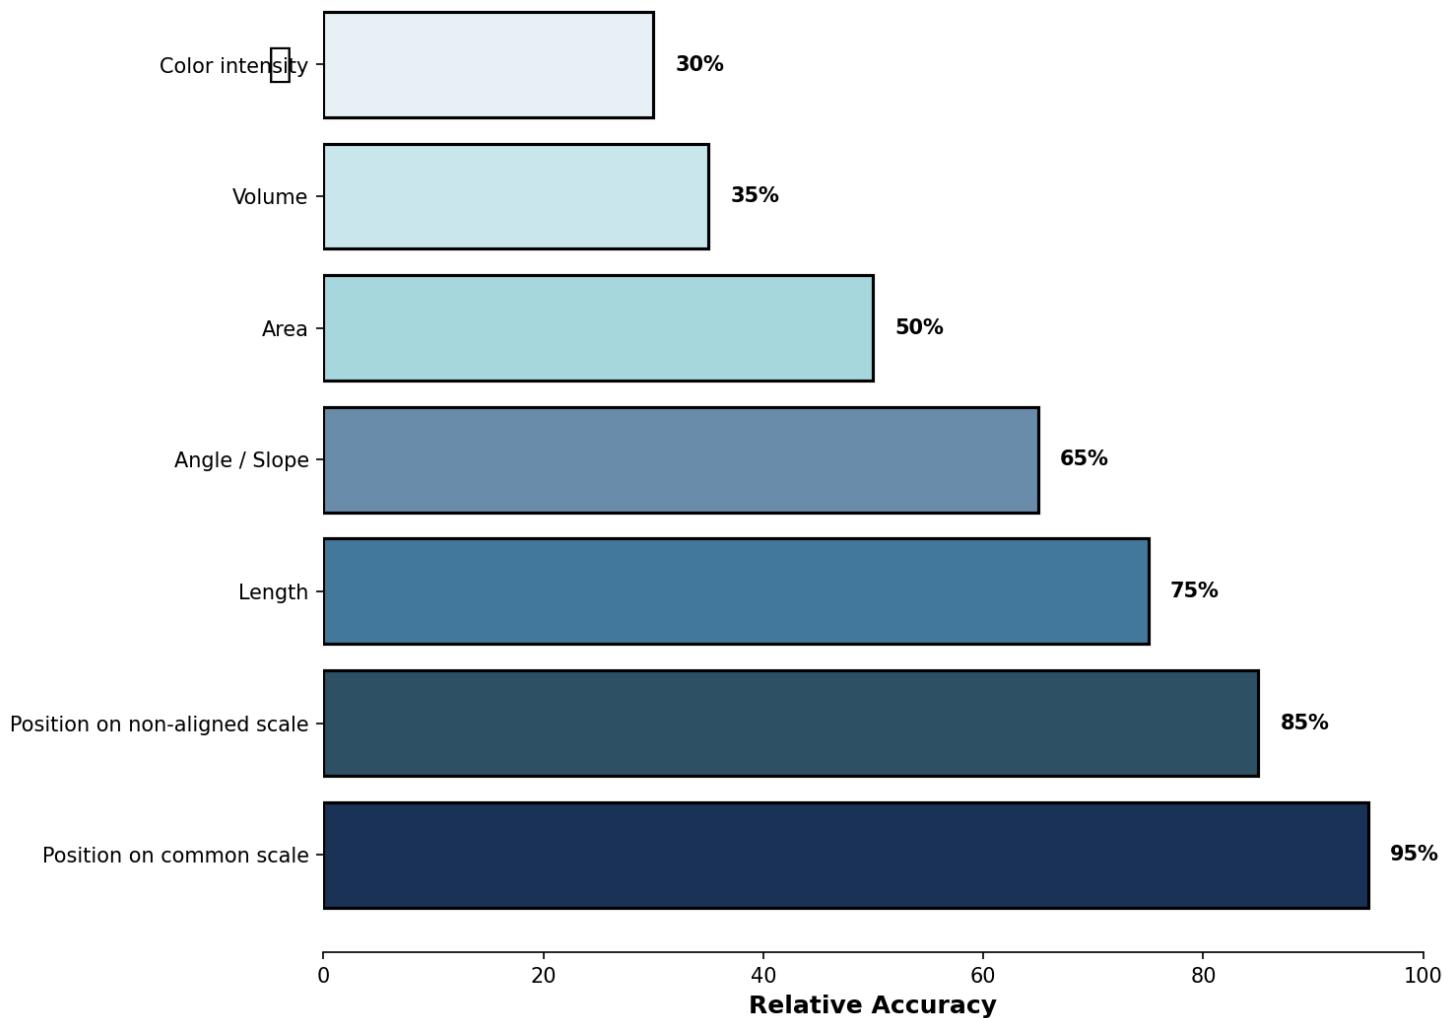
For Quantitative Data:

1. ★ Position on common scale (scatter plots, aligned bars)
2. Position on non-aligned scale (small multiples with different scales)
3. Length, direction, angle
4. Area
5. Volume, curvature
6. Shading, color saturation

Implications:

- Bar charts > Pie charts for comparison
- Dot plots > Bubble charts for precision
- Heatmaps good for patterns, not precise values

Visual Encoding Effectiveness (Most → Least Effective)



Why Color Is Weak for Quantitative Data

Color perception challenges:

- Non-linear perception (perceptual vs. actual lightness)
- Context-dependent (surrounding colors affect perception)
- Individual variation (color blindness affects ~8% of males)
- Poor for precise comparison

When to use color for quantitative data:

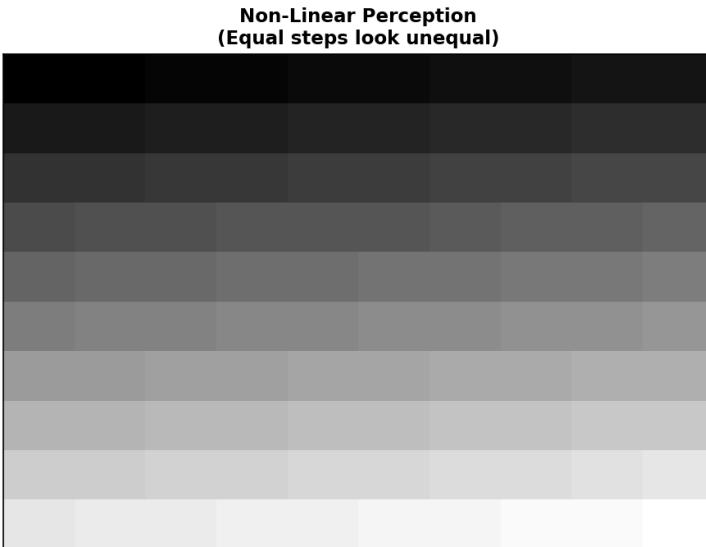
- Heatmaps (patterns, not precise values)
- Choropleth maps (general magnitude)
- Diverging scales (positive vs. negative)

When NOT to use color:

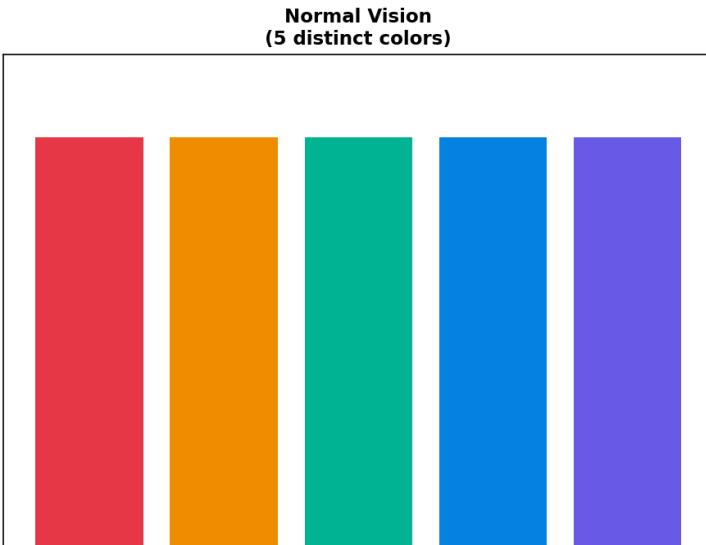
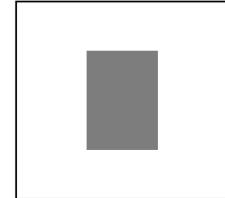
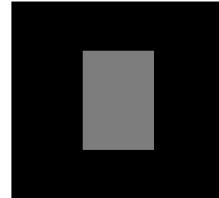
- Precise comparisons
- As the only encoding for critical information

Best practice: Use color for categories or as a secondary encoding, not primary for quantities.

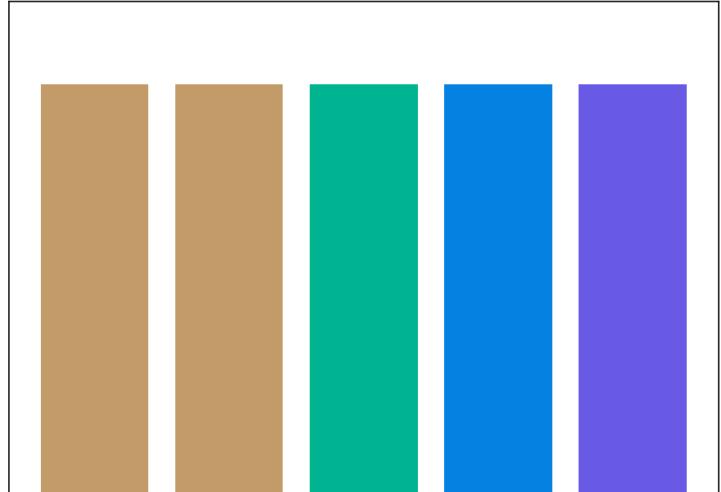
Why Color is Weak for Quantitative Data



Context Dependency
(Same gray looks different)



Deuteranopia (8% of males)
(Red & orange look similar)



Data × Task × Encoding Framework

Effective visualization requires matching all three:

$$\boxed{\text{Data Type}} + \boxed{\text{Task}} = \boxed{\text{Encoding Choice}}$$

Always ask:

1. What is the data type? (Nominal, Ordinal, Quantitative, Temporal)
2. What is the task? (Compare, Identify trend, Find outliers, etc.)
3. Which encoding best supports this?

Example:

- Data: Quarterly revenue (Quantitative + Temporal)
- Task: Show trend over time
- Encoding: Line chart with position on aligned scales

PART 3: GRAMMAR OF GRAPHICS FRAMEWORK

Slides 34-44

Common Encoding Violations

Mistakes that violate perception principles:

- ✗ Pie charts with too many slices (angle is hard to compare)
- ✗ 3D charts (perspective distorts perception)
- ✗ Dual-axis charts (misleading scale comparison)
- ✗ Rainbow color scales (non-perceptual, creates false boundaries)
- ✗ Truncated y-axes (exaggerates differences)
- ✗ Bubble charts for precise comparison (area is hard to judge)

Each of these increases cognitive load and decreases accuracy.

From Chart Types to Compositional Systems

Traditional approach:

- Think in fixed chart types (bar, line, pie, scatter)
- Pick a template
- Limited flexibility

Modern approach:

- Think in composable elements
- Build visualizations from components
- Infinite combinations possible

This is the shift from "choosing a chart" to "composing a visualization"

Enter: Grammar of Graphics

Introduction to Grammar of Graphics

Just as language has grammar for constructing sentences, visualization has grammar for constructing graphics

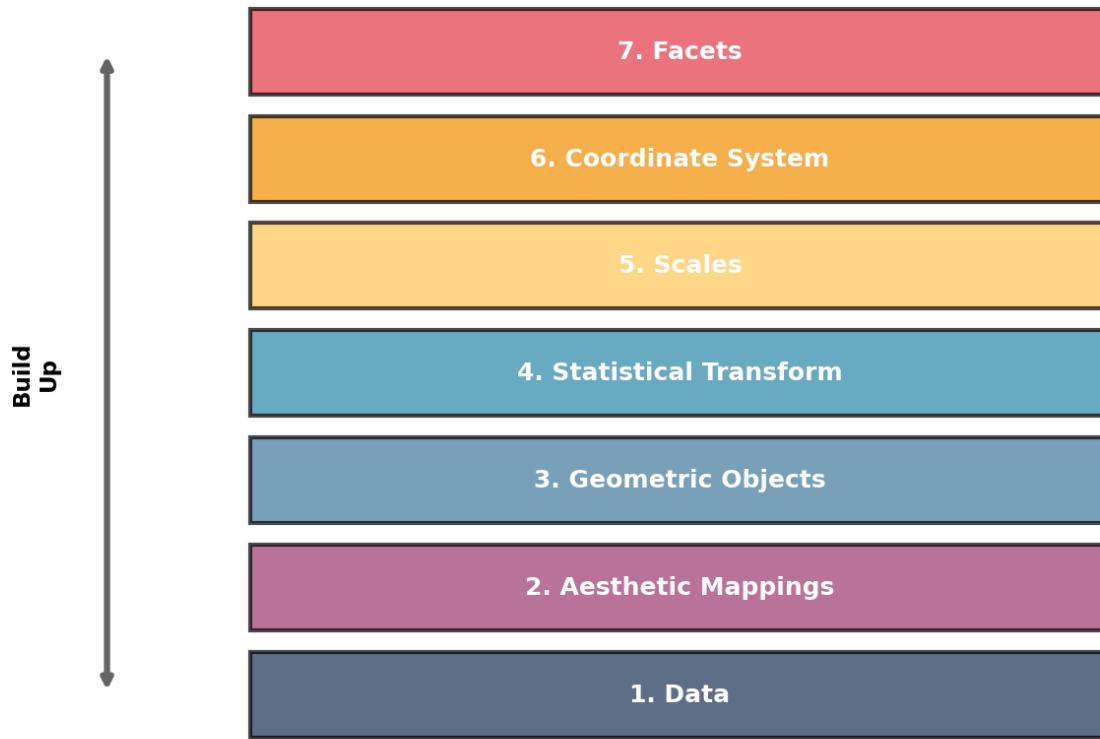
Developed by: Leland Wilkinson (1999), refined by Hadley Wickham (ggplot2)

Key insight: Visualizations are not types, but mappings from data to visual properties

Benefits:

- ⓘ Systematic approach to visualization design
- ⓘ Clear separation of concerns
- ⓘ More expressive than chart templates
- ⓘ Foundation for modern tools (ggplot2, Altair, Plotly)

Grammar of Graphics: Layered Structure



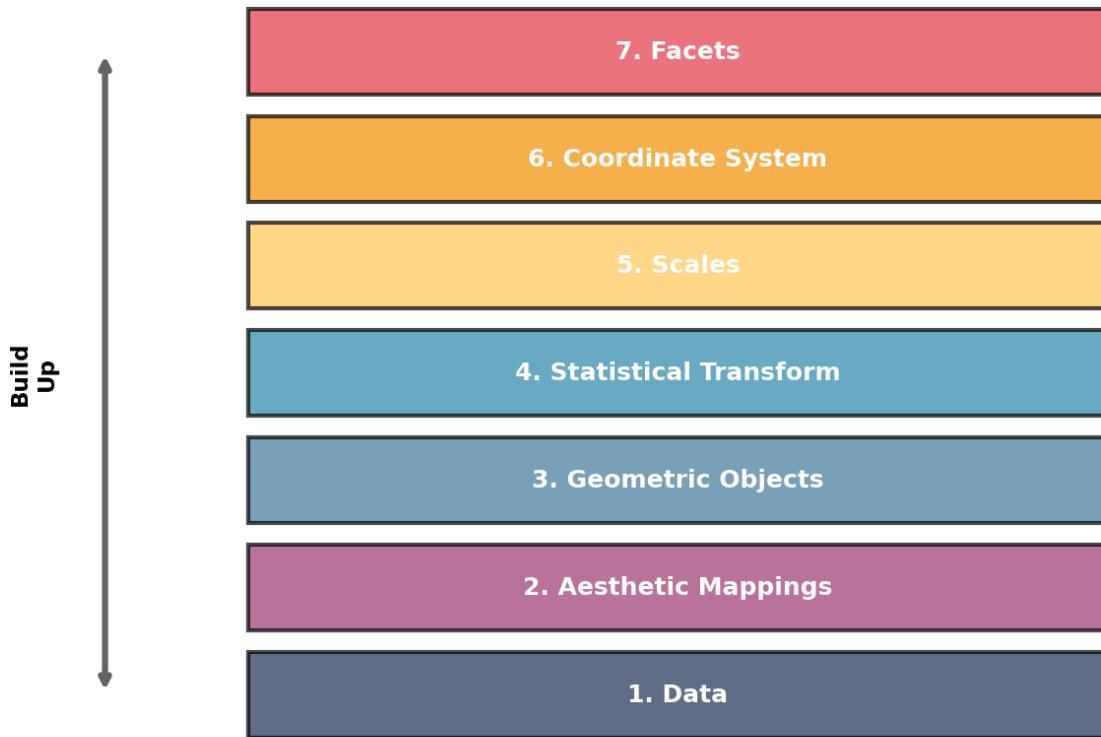
The Layered Grammar: Core Components

A visualization is built from layers, each with:

1. Data - The dataset to visualize
2. Aesthetic Mappings (`aes`) - Map data variables to visual properties
3. Geometric Objects (`geom`) - The visual marks (points, lines, bars)
4. Statistical Transformations (`stat`) - Data aggregations/computations
5. Scales - Map data ranges to visual ranges
6. Coordinate Systems - The space in which data is plotted
7. Facets - Subplots for different data subsets

Each component is independent and composable!

Grammar of Graphics: Layered Structure



Component 1: Data

The foundation of any visualization

Requirements:

- Structured format (typically tabular/dataframe)
- Each row = observation
- Each column = variable
- Can be raw or transformed

Example data structure:

Month	Product	Sales	Region
Jan	Widget A	1500	North
Jan	Widget B	2300	North
Feb	Widget A	1800	North
...			

Key insight: Same data + different mappings = different visualizations

Component 2: Aesthetic Mappings

Map data variables to visual properties

Common aesthetics:

- x , y - Position
- color / fill - Color of points/areas
- size - Size of points/text
- shape - Shape of points
- alpha - Transparency
- linetype - Type of line (solid, dashed)

Example:

```
aes(x = Month, y = Sales, color = Product, size = Region)
```

Critical: Aesthetics define *what* data is shown and *how* it maps to visual properties

Component 3: Geometric Objects (Marks)

The visual representation of data

Common geoms:

- point - Scatter plot points
- line - Connected lines
- bar / col - Bars/columns
- boxplot - Box-and-whisker plots
- histogram - Frequency distributions
- area - Filled areas
- tile / raster - Heatmap cells
- text - Labels

Same data + mapping, different geom = different chart:

- geom_point() → scatter plot
- geom_line() → line chart
- geom_bar() → bar chart

Component 4: Statistical Transformations

Compute derived values from data

Common transformations:

- identity - Use data as-is (default for scatter plots)
- count - Count observations (histograms)
- bin - Bin continuous data
- smooth - Fit smoothing line
- boxplot - Compute quartiles
- density - Compute density estimation

Example:

- Raw data: 1000 individual sales transactions
- After stat_count : Grouped counts per month
- After stat_smooth : Trend line with confidence interval

Often implicit: Bar chart automatically counts, histogram automatically bins

Component 5: Scales

Control how data values map to visual values

Examples:

- **Position scales:** Linear, logarithmic, date/time
- **Color scales:** Sequential, diverging, categorical palettes
- **Size scales:** Linear, square root (for area)

Key functions:

- Define domain (data range)
- Define range (visual range)
- Handle transformations
- Create legends/axes

Example:

```
Sales: $0 - $1M → Y-axis: 0px - 400px  
Product: A, B, C → Colors: blue, red, green
```

Component 6: Coordinate Systems

The space in which data is plotted

Common coordinate systems:

- **Cartesian** (x, y) - Most common, rectangular
- **Polar** - Circular (pie charts, radar charts)
- **Geographic** - Map projections
- **Transformed** - Log scale, square root

Impact:

- Same data, different coordinate system = dramatically different interpretation
- Cartesian bar chart vs. Polar (coxcomb) chart

Design tip: Don't use polar coordinates unless circular relationships are meaningful!

Component 7: Faceting (Small Multiples)

Create subplots for data subsets

Power of small multiples:

- Easy comparison across groups
- Reduces overplotting
- Maintains consistent scales
- Reduces cognitive load vs. overlaying

Types:

- **Grid:** Rows × Columns (e.g., by Region × Product)
- **Wrap:** Single variable, wraps to fit space

Example: Sales trends for each product category in separate panels

 Small Multiples

PART 4: PYTHON VISUALIZATION IMPLEMENTATION

Slides 45-82

Grammar and Cognitive Load

How Grammar of Graphics reduces cognitive load:

- ✓ Separation of concerns - Each component has one job
- ✓ Consistent structure - Same logic across all visualizations
- ✓ Explicit mappings - No hidden assumptions
- ✓ Composability - Build complex from simple
- ✓ Systematic thinking - Framework guides design decisions

Result: Less accidental complexity, more intentional design

The grammar makes design decisions explicit, reducing extraneous load.

Transition to Python: Implementing the Grammar

Now we'll see how to implement these concepts in Python

Python visualization ecosystem:

-  Matplotlib - Low-level, full control, foundation for others
-  Seaborn - Statistical graphics, built on Matplotlib
-  Plotly - Interactive visualizations
-  Altair - Declarative, pure Grammar of Graphics implementation

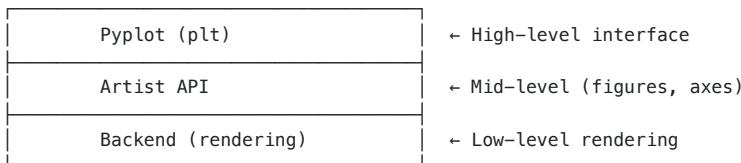
Today's focus: Matplotlib + Seaborn

Why?

- Industry standard
- Flexible and powerful
- Foundation for understanding other libraries

Python Visualization Ecosystem Overview

Matplotlib Architecture:



Seaborn sits on top:

- Provides high-level statistical plots
- Better default aesthetics
- Grammar-inspired interface
- Built on Matplotlib

Learning path: Start with pyplot, understand axes/figures, then use Seaborn for statistical plots.

Matplotlib: Anatomy of a Figure

Every Matplotlib visualization has two main components:

```
import matplotlib.pyplot as plt
```

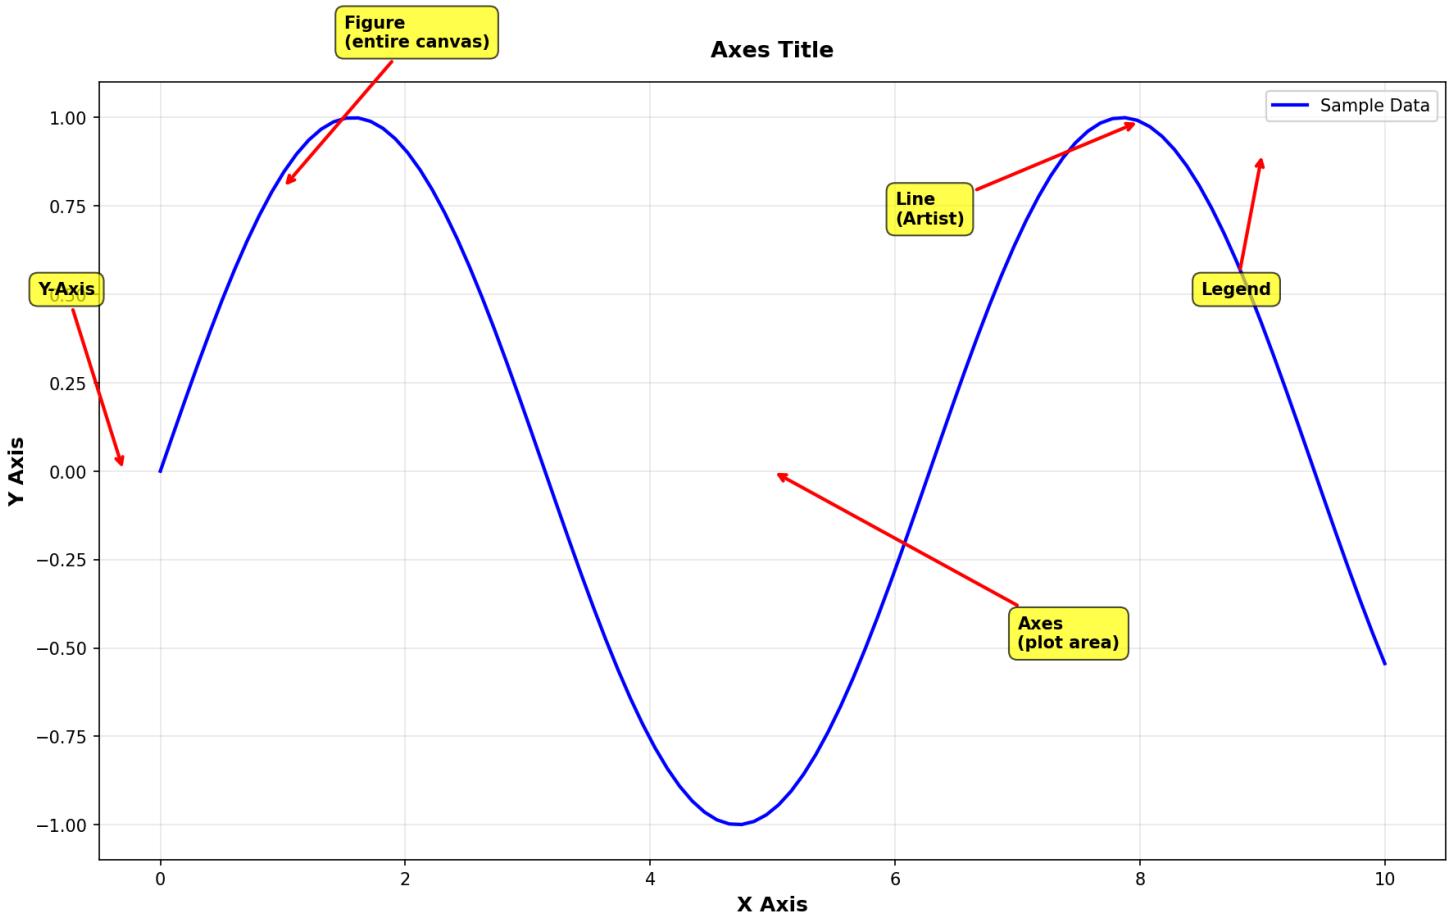
```
# Create figure and axes
fig, ax = plt.subplots() # Returns Figure and Axes objects

# fig = Canvas (the whole image)
# ax = Plot area (where data is drawn)
```

Key concepts:

- **Figure:** The entire canvas (can contain multiple plots)
- **Axes:** Individual plot(s) within the figure
- **Axis:** The x and y number lines
- **Artists:** Everything you see (lines, text, ticks, etc.)

Matplotlib Anatomy



Creating Your First Plot

Basic bar chart in Matplotlib:

```
import matplotlib.pyplot as plt
import numpy as np

# Data
categories = ['Product A', 'Product B', 'Product C', 'Product D']
sales = [2500, 3200, 2800, 3500]

# Create plot
fig, ax = plt.subplots(figsize=(8, 5))
ax.bar(categories, sales, color='steelblue')

# Customize
ax.set_xlabel('Product Category', fontsize=12)
ax.set_ylabel('Sales ($)', fontsize=12)
ax.set_title('Q1 2026 Sales by Product', fontsize=14, fontweight='bold')
```

```

# Add value labels on bars
for i, v in enumerate(sales):
    ax.text(i, v + 50, str(v), ha='center', fontweight='bold')

plt.tight_layout()
plt.show()

```

Output: Clean bar chart with direct labeling (reducing cognitive load!)

Customizing Plots: Colors, Labels, Styles

Applying perception principles:

```

import matplotlib.pyplot as plt

# Create data
months = ['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun']
sales_2025 = [45000, 48000, 52000, 49000, 55000, 58000]
sales_2026 = [47000, 51000, 56000, 54000, 60000, 65000]

# Create plot with good practices
fig, ax = plt.subplots(figsize=(10, 6))

# Use position (best encoding for quantitative comparison)
ax.plot(months, sales_2025, marker='o', linewidth=2,
        label='2025', color='#2E86AB') # Distinct colors
ax.plot(months, sales_2026, marker='s', linewidth=2,
        label='2026', color='#A23B72')

# Clear labels (reduce extraneous load)
ax.set_xlabel('Month', fontsize=11)
ax.set_ylabel('Sales ($)', fontsize=11)
ax.set_title('Sales Comparison: 2025 vs 2026',
             fontsize=13, fontweight='bold', pad=15)

# Legend with clear positioning
ax.legend(frameon=True, loc='upper left', fontsize=10)

# Remove top and right spines (reduce clutter)
ax.spines['top'].set_visible(False)
ax.spines['right'].set_visible(False)

# Add grid for easier reading (but subtle)
ax.grid(True, alpha=0.3, linestyle='--')

plt.tight_layout()
plt.show()

```

Matplotlib: Common Plot Types

Quick reference for basic plots:

```

import matplotlib.pyplot as plt
import numpy as np

# Sample data
x = np.arange(10)
y1 = np.random.rand(10) * 100
y2 = np.random.rand(10) * 100

fig, axes = plt.subplots(2, 3, figsize=(12, 8))

# Scatter plot
axes[0,0].scatter(y1, y2, alpha=0.6, s=100)
axes[0,0].set_title('Scatter Plot')

# Line plot

```

```

axes[0,1].plot(x, y1, marker='o')
axes[0,1].set_title('Line Plot')

# Bar chart
axes[0,2].bar(x, y1, color='coral')
axes[0,2].set_title('Bar Chart')

# Histogram
axes[1,0].hist(y1, bins=10, edgecolor='black', alpha=0.7)
axes[1,0].set_title('Histogram')

# Box plot
axes[1,1].boxplot([y1, y2], labels=['Group A', 'Group B'])
axes[1,1].set_title('Box Plot')

# Heatmap (using imshow)
data = np.random.rand(10, 10)
im = axes[1,2].imshow(data, cmap='YlOrRd')
axes[1,2].set_title('Heatmap')
plt.colorbar(im, ax=axes[1,2])

plt.tight_layout()
plt.show()

```

Key insight: Same pyplot interface, different geom!

Introduction to Seaborn

Seaborn = Statistical Data Visualization

Built on Matplotlib, but provides:

- Better default aesthetics
- Grammar-inspired interface
- Statistical transformations built-in
- Works directly with pandas DataFrames
- Themes and color palettes

Installation:

```
pip install seaborn
```

Import convention:

```
import seaborn as sns
import matplotlib.pyplot as plt
```

Philosophy: Make complex statistical plots simple, apply perception principles by default.

Seaborn: Working with DataFrames

Seaborn loves pandas DataFrames:

```

import seaborn as sns
import pandas as pd

# Create sample data
data = pd.DataFrame({
    'month': ['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun'] * 2,
    'sales': [2500, 2700, 3000, 2900, 3200, 3400,
              2600, 2800, 3100, 3000, 3300, 3500],
    'region': ['North']*6 + ['South']*6
})

# Create plot using column names (Grammar-like!)
sns.lineplot(data=data, x='month', y='sales', hue='region',

```

```
marker='o', linewidth=2)
plt.title('Sales by Region Over Time')
plt.show()
```

Notice: Direct mapping from data columns to aesthetics!

Seaborn Plot Types

Common Seaborn plots mapped to tasks:

Distribution:

- `histplot()` - Histograms with KDE
- `kdeplot()` - Kernel density estimation
- `boxplot()` - Box-and-whisker
- `violinplot()` - Violin plot (box + density)

Relationships:

- `scatterplot()` - Scatter with regression
- `lineplot()` - Line with confidence intervals
- `regplot()` - Scatter + regression line

Categories:

- `barplot()` - Bars with confidence intervals
- `countplot()` - Count plot
- `pointplot()` - Point estimates with error bars

Seaborn Example: Distribution Plot

Visualizing distributions with built-in statistics:

```
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt

# Generate sample data
np.random.seed(42)
data1 = np.random.normal(100, 15, 500)
data2 = np.random.normal(120, 20, 500)

# Create figure
fig, axes = plt.subplots(1, 2, figsize=(12, 5))

# Histogram with KDE
sns.histplot(data1, kde=True, color='steelblue', ax=axes[0])
axes[0].set_title('Sales Distribution - Product A')
axes[0].set_xlabel('Daily Sales ($)')

# Violin plot comparing two groups
import pandas as pd
df = pd.DataFrame({
    'Sales': np.concatenate([data1, data2]),
    'Product': ['A']*500 + ['B']*500
})

sns.violinplot(data=df, x='Product', y='Sales', ax=axes[1])
axes[1].set_title('Sales Distribution Comparison')

plt.tight_layout()
plt.show()
```

Automatic statistical transformations!

Seaborn: Faceting (Small Multiples)

Create small multiples with FacetGrid:

```
import seaborn as sns
import pandas as pd

# Load sample data
tips = sns.load_dataset('tips')

# Create FacetGrid (Component 7: Faceting!)
g = sns.FacetGrid(tips, col='time', row='smoker', height=4)
g.map(sns.scatterplot, 'total_bill', 'tip', alpha=0.7)
g.add_legend()
g.set_titles(col_template="{col_name}", row_template="{row_name}")
g.set_axis_labels("Total Bill ($)", "Tip ($)")

plt.show()
```

Result: 4 panels (Lunch/Dinner × Smoker/Non-smoker)

Grammar in action:

- Data: tips dataset
- Aesthetics: x=total_bill, y=tip
- Geom: scatter
- Facets: time × smoker

Color Palettes in Python: Perception-Aware Choices

Choosing the right color palette matters!

```
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

# Sample data
data = np.random.randn(100, 4)

fig, axes = plt.subplots(1, 3, figsize=(15, 4))

# ❌ BAD: Rainbow (perceptually non-uniform, creates false boundaries)
axes[0].imshow(data, cmap='jet', aspect='auto')
axes[0].set_title('❌ Rainbow (Jet) - Avoid!')

# ✅ GOOD: Sequential (perceptually uniform)
axes[1].imshow(data, cmap='viridis', aspect='auto')
axes[1].set_title('✅ Viridis - Sequential Data')

# ✅ GOOD: Diverging (for data with meaningful midpoint)
axes[2].imshow(data, cmap='RdBu_r', aspect='auto')
axes[2].set_title('✅ Red-Blue - Diverging Data')

plt.tight_layout()
plt.show()
```

Seaborn color palettes:

- Categorical: 'Set2', 'tab10', 'Paired'
- Sequential: 'Blues', 'Greens', 'rocket'
- Diverging: 'RdBu', 'coolwarm', 'vlag'

Applying Preattentive Features in Code

Using color for visual pop-out:

```

import matplotlib.pyplot as plt
import numpy as np

# Data
categories = ['Q1', 'Q2', 'Q3', 'Q4']
sales_2025 = [450, 480, 520, 490]
sales_2026 = [470, 510, 560, 540]

fig, axes = plt.subplots(1, 2, figsize=(12, 5))

# WITHOUT preattentive highlighting
axes[0].bar(categories, sales_2025, color='steelblue', alpha=0.7)
axes[0].set_title('Without Visual Emphasis')
axes[0].set_ylabel('Sales (K$)')

# WITH preattentive highlighting (Q3 is best quarter)
colors = ['#CCCCCC', '#CCCCCC', '#E63946', '#CCCCCC'] # Pop-out!
axes[1].bar(categories, sales_2026, color=colors, alpha=0.9)
axes[1].set_title('With Visual Emphasis - Q3 Highlighted')
axes[1].set_ylabel('Sales (K$)')

plt.tight_layout()
plt.show()

```

Result: Q3 immediately pops out (preattentive color feature)

Reducing Clutter: Before and After

Applying "reduce extraneous load" principle:

```

import matplotlib.pyplot as plt
import numpy as np

months = ['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun']
sales = [2500, 2700, 3000, 2900, 3200, 3400]

fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# ❌ CLUTTERED (high extraneous load)
axes[0].plot(months, sales, marker='o', linewidth=3, color='red',
             markersize=12)
axes[0].set_title('BEFORE: High Extraneous Load', fontsize=14,
                  fontweight='bold', color='darkred')
axes[0].set_facecolor('#FFEEEE') # Distracting background
axes[0].grid(True, linewidth=2, color='blue') # Heavy grid
axes[0].spines['top'].set_linewidth(3)
axes[0].spines['right'].set_linewidth(3)
axes[0].spines['bottom'].set_linewidth(3)
axes[0].spines['left'].set_linewidth(3)

# ✅ CLEAN (low extraneous load)
axes[1].plot(months, sales, marker='o', linewidth=2, color="#2E86AB",
             markersize=7)
axes[1].set_title('AFTER: Low Extraneous Load', fontsize=13)
axes[1].set_xlabel('Month', fontsize=10)
axes[1].set_ylabel('Sales ($)', fontsize=10)
axes[1].spines['top'].set_visible(False) # Remove clutter
axes[1].spines['right'].set_visible(False)
axes[1].grid(True, alpha=0.2, linestyle='--') # Subtle grid

plt.tight_layout()
plt.show()

```

Direct Labeling vs. Legend

Reduce cognitive load by eliminating legend lookup:

```

import matplotlib.pyplot as plt

months = ['Jan', 'Feb', 'Mar', 'Apr']
product_a = [2500, 2700, 3000, 2900]
product_b = [2200, 2400, 2600, 2500]
product_c = [1800, 2000, 2200, 2100]

fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# WITH LEGEND (requires lookup - adds cognitive load)
axes[0].plot(months, product_a, marker='o', label='Product A')
axes[0].plot(months, product_b, marker='s', label='Product B')
axes[0].plot(months, product_c, marker='^', label='Product C')
axes[0].legend()
axes[0].set_title('With Legend (More Cognitive Load)')
axes[0].set_ylabel('Sales ($)')

# WITH DIRECT LABELS (immediate understanding)
axes[1].plot(months, product_a, marker='o', color="#2E86AB")
axes[1].plot(months, product_b, marker='s', color="#A23B72")
axes[1].plot(months, product_c, marker='^', color="#F18F01")
# Add direct labels
axes[1].text(3.05, product_a[-1], 'Product A', va='center', fontsize=10,
            color="#2E86AB", fontweight='bold')
axes[1].text(3.05, product_b[-1], 'Product B', va='center', fontsize=10,
            color="#A23B72", fontweight='bold')
axes[1].text(3.05, product_c[-1], 'Product C', va='center', fontsize=10,
            color="#F18F01", fontweight='bold')
axes[1].set_title('With Direct Labels (Less Cognitive Load)')
axes[1].set_ylabel('Sales ($)')
axes[1].set_xlim(-0.2, 3.8)

plt.tight_layout()
plt.show()

```

Principle: Direct labeling reduces the need to constantly refer back to legend.

Exercise: Bar Chart vs. Pie Chart

Implementing Cleveland & McGill's findings:

```

import matplotlib.pyplot as plt

# Data
categories = ['North', 'South', 'East', 'West', 'Central']
values = [25, 22, 18, 20, 15]

fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# PIE CHART (uses angle - weak encoding)
axes[0].pie(values, labels=categories, autopct='%1.1f%%', startangle=90)
axes[0].set_title('Pie Chart: Angle Encoding\n(Harder to Compare)')

# BAR CHART (uses position - strong encoding)
axes[1].barh(categories, values, color='steelblue')
axes[1].set_xlabel('Sales (%)')
axes[1].set_title('Bar Chart: Position Encoding\n(Easier to Compare)')
axes[1].spines['top'].set_visible(False)
axes[1].spines['right'].set_visible(False)

# Add value labels
for i, v in enumerate(values):
    axes[1].text(v + 0.5, i, str(v) + '%', va='center')

plt.tight_layout()
plt.show()

```

Question: Which makes it easier to see that North > South > West?

Answer: Bar chart (position is more accurate than angle)

Putting It All Together: Complete Example

Building a perceptually-optimized visualization from scratch:

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Sample data
data = pd.DataFrame({
    'Quarter': ['Q1', 'Q2', 'Q3', 'Q4'] * 3,
    'Revenue': [450, 480, 520, 490, 470, 510, 560, 540, 490, 530, 580, 560],
    'Year': ['2024']*4 + ['2025']*4 + ['2026']*4
})

# Apply all principles
sns.set_style("whitegrid", {'grid.linestyle': '--', 'grid.alpha': 0.3})
fig, ax = plt.subplots(figsize=(10, 6))

# Use position encoding (best for quantitative)
sns.lineplot(data=data, x='Quarter', y='Revenue', hue='Year',
              marker='o', linewidth=2.5, palette=['#CCCCCC', '#999999', '#E63946'],
              ax=ax)

# Clear, direct labels
ax.set_xlabel('Quarter', fontsize=12)
ax.set_ylabel('Revenue ($K)', fontsize=12)
ax.set_title('Quarterly Revenue Growth 2024–2026\n2026 Shows Strong Performance',
             fontsize=14, fontweight='bold', pad=20)

# Reduce extraneous load
ax.spines['top'].set_visible(False)
ax.spines['right'].set_visible(False)
ax.legend(title='Year', frameon=False, loc='upper left')

# Add annotation for insight
ax.annotate('Best Quarter', xy=('Q3', 580), xytext=('Q4', 595),
            arrowprops=dict(arrowstyle='->', color='#E63946', lw=2),
            fontsize=10, color='#E63946', fontweight='bold')

plt.tight_layout()
plt.show()
```

Principles applied: Position encoding, color for emphasis, reduced clutter, direct annotation

Common Mistakes to Avoid

Real-world examples of what NOT to do:

1. ✗ Too many colors - Rainbow schemes for categories
2. ✗ 3D without purpose - Distorts perception
3. ✗ Dual Y-axes - Misleading comparisons
4. ✗ Truncated axes - Exaggerates differences
5. ✗ Pie charts with 10+ slices - Impossible to compare
6. ✗ Busy backgrounds - Distracts from data
7. ✗ Tiny fonts - Can't read = won't use
8. ✗ No labels - Forces guessing

Remember: Every element should serve the data and the user's task!

Visualization Design Checklist

Before finalizing any visualization, ask:

1. Data & Task

- What data type am I showing?
- What task does the user need to perform?
- Is this the right chart for the task?

2. Perception

- Am I using the most effective encoding?
- Are preattentive features highlighting what matters?
- Can users see the pattern in < 5 seconds?

3. Cognitive Load

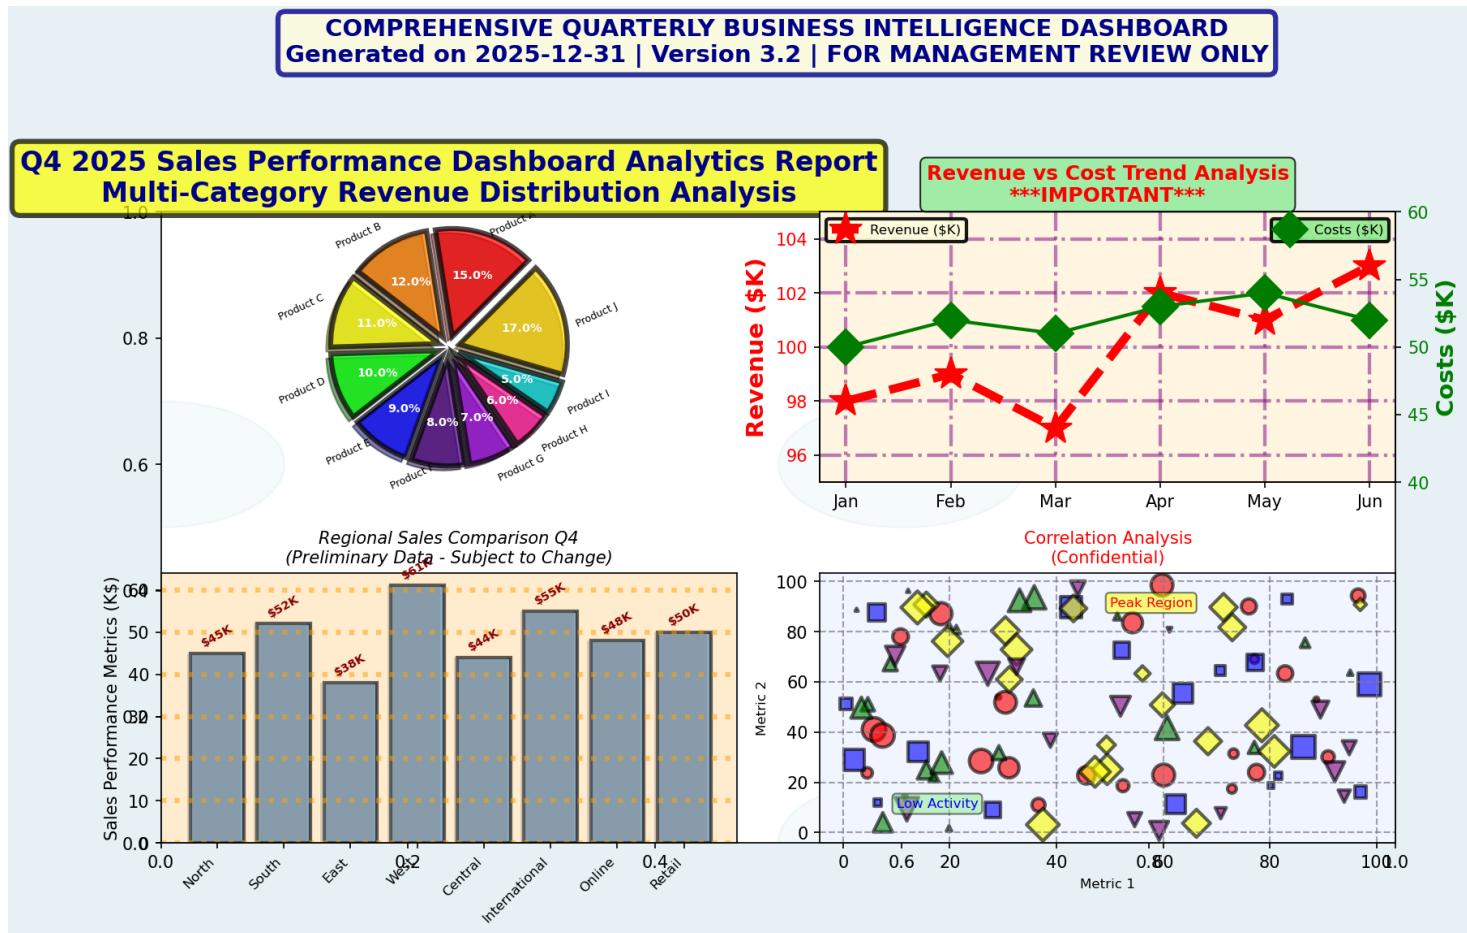
- Have I removed all chart junk?
- Is the visual hierarchy clear?
- Are labels direct and clear?

4. Accessibility

- Is it colorblind-friendly?
- Are fonts large enough?
- Does it work in grayscale?

Exercise 1: Critique This Visualization

Task: Identify all perception and cognitive load issues in this dashboard



Issues to find:

1. Encoding problems - What visual encodings are inappropriate?
2. Cognitive load violations - What adds unnecessary complexity?
3. Gestalt principle violations - What grouping issues exist?
4. Color usage problems - Where is color misused?
5. Missing context - What information is unclear or missing?

Hint: Look for:

- Pie charts with too many slices
- Rainbow/non-perceptual color scales
- Truncated axes
- Dual y-axes with different scales
- Chart junk (decorative elements)
- Poor labeling and tiny text
- Inconsistent styling
- Cluttered layouts

Work in pairs: 5 minutes to list all issues you can find

Expected findings: This visualization violates at least 10 principles we learned today!

Exercise 2: Redesign Challenge

Given: Bad dashboard with high cognitive load

Your task:

1. List all sources of extraneous load
2. Propose a redesigned layout
3. Justify your encoding choices
4. Apply Gestalt principles

Deliverable: Sketch or code the improved version

Time: 15 minutes

Exercise 3: Build a Grammar-Based Plot

Task: Create a visualization using Grammar of Graphics thinking

Dataset: Sales by product, region, and quarter

Requirements:

1. Define your data mappings (aesthetics)
2. Choose appropriate geom
3. Consider statistical transformation
4. Use faceting if appropriate
5. Apply perception principles

Implement in Python (Matplotlib or Seaborn)

Exercise 4: Color Palette Selection

Scenario: You need to visualize temperature data from -10°C to +40°C

Questions:

1. What type of color palette should you use?
2. Why is rainbow/jet inappropriate?
3. What's a good alternative?
4. How do you handle the zero point?

Implement:

```
import matplotlib.pyplot as plt
import numpy as np

# Your code here to create heatmap with appropriate palette
```

Exercise 5: Preattentive Highlighting

Task: Create a sales report where Q4 must pop out

Data: Quarterly sales for 5 products

Requirements:

- Use preattentive features (color, size, position)
- Make Q4 immediately visible
- Minimize cognitive load
- Include direct labels

Code it: Use Matplotlib or Seaborn

Real-World Application: Dashboard Design

Scenario: Executive wants dashboard for:

- Monthly revenue trends
- Top 5 products
- Regional performance
- YoY comparison

Your design must:

1. Apply visual hierarchy (Gestalt)
2. Use appropriate encodings
3. Minimize cognitive load
4. Enable 5-second insights

Discussion: What would you include? What would you exclude?

Connection to Data Science Workflows

Where visualization fits in your projects:

Data Collection → Cleaning → EDA → Modeling → Evaluation → Communication
↓ ↓ ↓ ↓
Visualization is critical at every stage!

Roles of visualization:

- **EDA:** Understand patterns, find outliers
- **Model Selection:** Compare performance
- **Debugging:** Identify issues
- **Communication:** Share insights
- **Decision Support:** Enable action

Remember: Good visualization = Better models = Better decisions